KARTHIK PAPER WALLET

Track Record Moonshot Strategy

One rule, stated once, and nothing else: **$1,000 of capital. $10 into every new Track Record token. Sell all of it at 1.25×. No stop loss. No time exit.** Otherwise hold until the token is genuinely dead.

STARTING CAPITAL	TRADE SIZE	PROFIT TARGET	STOP LOSS
$1,000.00	**$10.00**	**1.25×**	**NONE**
copied onto the wallet row at activation	fixed — not a maximum, not a fraction	sell 100%. One number, one exit	a loser may go to $0.00

VERIFIED IN PRODUCTION	VALUE
Wallet ID	`a0da0193-4fea-4ac1-891b-44abc897dc84`
Name	Karthik LIVE · TRADING
Activated	`2026-08-22 10:57:12.738872Z`
Starting capital	`1000.0000`
Trade size	`10.0000`
Take profit multiple	`1.2500`
Entries paused	`KARTHIK_ENTRIES_PAUSED=false` — entries are running
Review cadence	One pass per minute, `crontab(minute="*")`
Execution	`PAPER_EXECUTION_MODEL=jupiter` · 50 bps slippage

SOURCE OF OPPORTUNITIES

Every row in `radar_tokens` — the Track Record — whose `first_detected_at` is **later than the activation instant**. Nothing before it is ever eligible; nothing else is consulted. No score, rank or category filter is re-applied here, so Karthik's universe cannot drift from what the Track Record page shows.

WHY THERE IS NO STOP LOSS

A stop loss is a rule about price. “Goes to zero” is a claim about executability. A token at 0.05× is held exactly like a token at 1.10× is held; only the provider's own report that the pool has no meaningful liquidity left ends the trade.

WHY A PRINT IS NOT A FILL

Reaching 1.25× is **necessary but never sufficient**. A drained pool can print any number, and a rule that read the print would turn a rug into a winning trade. The position only closes when a router will quote the whole position — and the proceeds are that quote's, never `1.25 × cost`.

Every expectation in the brief was checked against production, and all five hold. $1,000 starting capital ✓ · $10 trade size ✓ · new Track Record tokens only ✓ · 1.25× profit target ✓ · no stop loss ✓. What the brief did not anticipate is on page 7: a target hit is **not** guaranteed to be a profit, and two of them so far were not.

02 Every decision Karthik makes

Corrected against the implementation. Two things the conceptual sketch did not have: **exits are settled before entries** and never consult cash, and a skip is **permanent** — a token Karthik could not afford or could not price is never queued and never bought later.

STEP 1 · SETTLE EXITS (runs first, every pass, never consults cash)

STEP 2 · DECIDE NEW ADMISSIONS (independent of step 1)

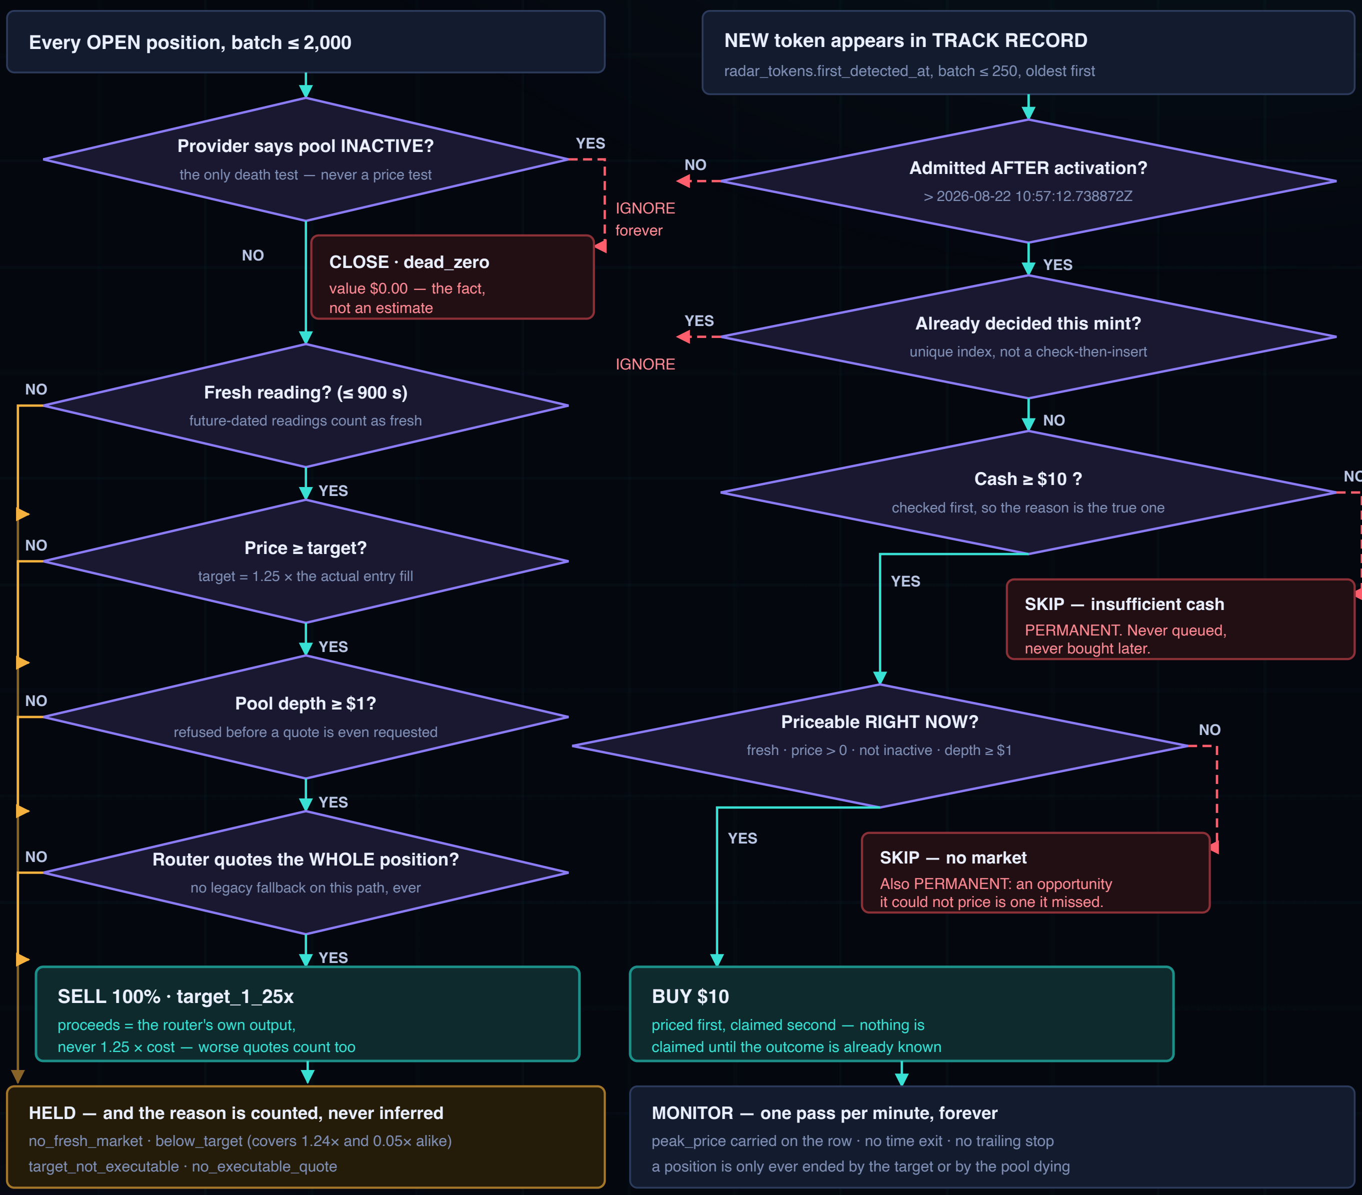

DUPLICATE PROTECTION

Two unique indexes, not application checks: `uq_karthik_opportunities_wallet_mint` and `uq_karthik_positions_wallet_mint`. A mint gets **one decision and at most one position, ever**. Two concurrent passes cost one wasted HTTP call and cannot cost a duplicate. A per-pass advisory lock (`"MEME"/"KRTH"`) is a convenience on top, not the guarantee.

THE TWO PRICES

Karthik decides from a stored observation and executes against a live router quote. The observation says **whether**; the quote says **at what**. An entry seen twenty minutes late is bought at the twenty-minutes-late price, because that is the price Karthik could have paid.

THE PAUSE SWITCH

`KARTHIK_ENTRIES_PAUSED` is **false** in production — entries are running. When set, it stops entries only: it is read inside step 2, which runs after exits have already settled, so it cannot strand an open book.

03 What $10 does, three ways

TRADE A · TARGET HIT

Invest	$10.00
Reaches	1.25× the entry fill
Router quotes	the whole position
Sell	**≈ $12.50**
Gross profit	**≈ +$2.50**

Nominal. What actually lands is the router's number.

TRADE B · NOT THERE YET

Invest	$10.00
Reaches	1.10×
Target	not met
Action	**keep holding**
Realised	**$0.00 — nothing**

Held as `below_target`. No stop, no clock. The $10 stays locked up.

TRADE C · THE RUG

Invest	$10.00
Pool	reported `inactive`
Sellable	nothing
Value	**$0.00**
Loss	**−$10.00**

The whole stake. There is no rule that can make this smaller.

THE BREAK-EVEN ARITHMETIC

+$2.50 +$2.50 +$2.50 +$2.50 = −$10.00

FOUR clean +25% winners · four separate $10 trades that each reached target

ONE token that went to zero

...and that is only to get back to flat.

Before any cost.

REQUIRED WIN RATE, FROM THE NOMINAL RULE

If a winner returns $2.50 and a total loss costs $10, the break-even hit rate p solves `2.50p = 10(1−p)`:

$$p = 80\%$$

Four out of every five positions must reach 1.25× just to break even — and that assumes every non-winner is a total loss and every winner returns exactly the nominal amount. Neither is true.

REQUIRED WIN RATE, FROM KARTHIK'S OWN FILLS

Production's 53 target hits returned **$14.085 on average** — above the nominal $12.50, because a few sold well past the target. Using the real average, break-even is `10 / 14.085`:

$$p = 71.0\%$$

But the average is carried by one $68.73 outlier. At the **median** of $12.837 the requirement is **77.9%**. **Karthik's actual hit rate is 53 / 80 = 66.3%** — short of every one of these thresholds, which is why realised P&L is negative.

The execution-cost model, measured rather than assumed. Karthik quotes every fill through Jupiter (`PAPER_EXECUTION_MODEL=jupiter`, 50 bps slippage tolerance); all 131 entries were priced by a real route. Across those entries the **average entry price impact was 10.24%** and the average exit impact **0.87%**; the platform fee recorded was **$0.00**. Costs are taken out of the tokens received, never recorded beside a full-size fill — cash always moves by exactly $10. Where no route is available the entry falls back to a deterministic model charging **30 bps per side** plus exact constant-product impact; if that exceeds $10 there is no entry at all. **There is no fallback on the exit path, deliberately** — the worst case there would be a rug booked as a 25% win.

04 How $1,000 spreads across $10 positions

Every position is exactly $10, so the capital picture is simple arithmetic: the wallet has **100 slots**, and each one is 1% of the book. Cash returns only when a position hits 1.25× — a dead token returns nothing and frees nothing, because it is closed at $0.00.

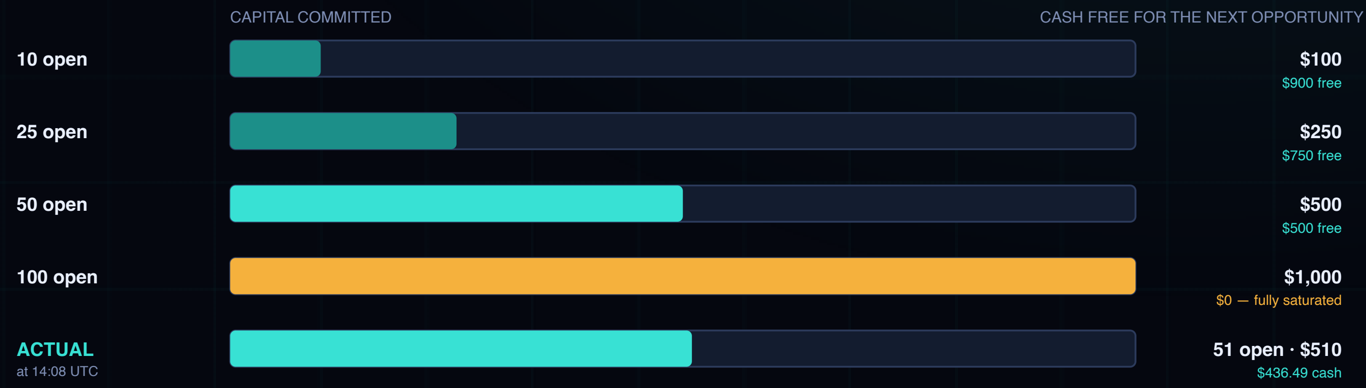

Cash is not $1,000 – allocated. It is `starting capital – everything ever committed + everything ever returned`. Karthik has committed **$1,310** across 131 entries and had **$746.49** returned, so cash is **$436.49**, not $490. The $53.51 difference is realised loss — it is gone, and it never comes back into the slot count.

WHAT $10 SIZING GENUINELY BUYS

- A single rug costs **1% of the book**, not 10%. At $100 per trade the same rug would cost ten times as much.
- **Far more opportunities captured.** Karthik decided **131 of 131** admissions since activation — a 100% capture rate — precisely because it could still afford every one of them.
- The **sample size** needed to judge the strategy arrives far sooner. 80 closed trades in three hours is a measurement; eight would have been an anecdote.

WHAT IT DOES NOT BUY

- **It does not create positive expectancy.** Smaller positions scale every outcome down together — the winners shrink by exactly the same factor as the losses. The ratio is untouched.
- The break-even hit rate is **identical at $10 and at $100**: 80% nominally. Position size cannot move it; only the win rate or the size of the wins can.
- More positions means **more exposure to the same risk**, not less. 100 open $10 positions is still $1,000 exposed to the same market.

ADMISSIONS SINCE ACTIVATION	DECIDED	SKIPPED FOR CASH	SKIPPED, NO MARKET
131	**131**	**0**	**0**
Track Record rows after 10:57 UTC	100% capture · all `entered`	has not yet run out	every admission was priceable

05 Agent Karthik, and what it is allowed to do

Agent Karthik is not running in production. Its source (`app/karthik_ops`) exists on the host checkout at commit `11fe507`, but it is **absent from the running container image**, which was built at `bcc5550`. Production compose declares `volumes: !reset []` — there is no bind mount — so code on disk never reaches the running process without a rebuild. **It serves no routes and takes no action today.** Everything on this page describes what is built, clearly separated from what is deployed.

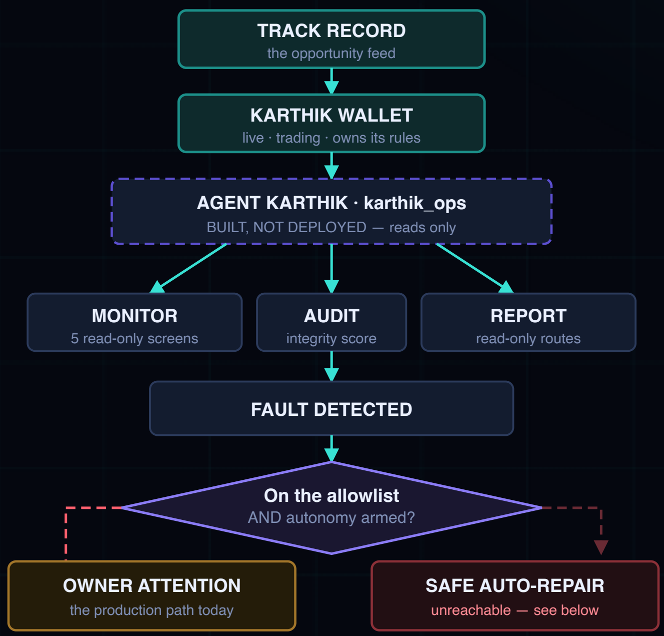

AUTONOMY MODE: OBSERVE-ONLY

Two modes exist in the code: `OBSERVE_ONLY` and `SAFE_AUTOREPAIR`. Anything unrecognised reads as `OBSERVE_ONLY`, and that is the production default.

No autonomous repairs are currently executed. This is true twice over: the module is not deployed, and the arming variable is unset. A known fault under `OBSERVE_ONLY` is **refused** and routed to the owner-approval queue instead.

WHAT IT IS FORBIDDEN TO TOUCH

Karthik is an operational layer. It does not decide what the wallet trades, when it enters, or when it exits — those are the wallet's own rules, and every one of them is on the `FORBIDDEN_STRATEGY_IDS` list it may not touch. Its API router carries **no POST, PUT, PATCH or DELETE**.

CORRECT WHILE UNBOUND

The module was written before the wallet existed. Every surface reports **NOT DESIGNATED** rather than zero, and the integrity score is `None` rather than 100 — no report claims a figure it could not read. Binding it is one variable, `KARTHIK_WALLET_STRATEGY_ID`, which is **not set anywhere in production**.

WHAT IS DEPLOYED AND WATCHING, TODAY

`hq_ops` — the platform-wide HQ layer — **is** in the running image and does run. It probes, classifies conditions, opens idempotent incidents and writes an audit row before it would act. Its own arming variable `HQ_AUTONOMY_ENABLED` is unset, so it too defaults to **false** and executes no repair. Record to date: **7 incidents, all resolved, none open**. HQ is not Karthik-specific — it watches the platform, and it does not know the Karthik wallet's rules.

06 The wallet, as measured

Read from the `karthik_*` tables and the production environment. Cells marked NOT SERVED are defined in `karthik_ops`, which is not in the running image — shown as unavailable, never as zero.

CAPITAL

WALLET EQUITY	AVAILABLE CASH	CAPITAL ALLOCATED	REALISED P&L
$946.49	$436.49	$510.00	–$53.51
cash + open book at cost	1,000 – 1,310 + 746.49	51 slots of 100	–5.35% of starting capital

POSITIONS

OPEN POSITIONS	CLOSED POSITIONS	TARGET HITS	DEAD / ZERO
51	80	53	27
no time limit on any of them	in the first 3 h 11 m	exit_reason target_1_25x	exit_reason dead_zero · $0.00 each

RATES

TARGET HIT RATE	RUG / DEAD RATE	BREAK-EVEN NEEDED	OPPORTUNITY CAPTURE
66.3%	33.8%	71–80%	100%
53 of 80 closed	27 of 80 closed	at realised avg → nominal	131 of 131 admissions decided

PIPELINE & INTEGRITY

SIGNAL	READING
Track Record → entry delay	mean **29.4 s** · min 0.7 s · max 119.6 s
Quote freshness rule	observations older than **900 s** are refused
Market feed lag	newest snapshot **30 s** old at verification
Execution confidence	**131 of 131** entries via `jupiter_route` — no fallbacks used
Worker health	worker, scheduler, scanner all **healthy**; enrichment ×2 up
Review cadence	every minute, own lock — never blocked by the other wallet

SIGNAL	READING
Experiment integrity	the activation instant is permanent by construction — a unique index makes a second wallet impossible
Duplicate protection	0 duplicates. 131 admissions → 131 decisions → 131 positions
Incidents (HQ, platform-wide)	**7 on record, all resolved, 0 open**
Owner attention queue	NOT SERVED — `karthik_ops` is not in the running image
Integrity score	NOT SERVED — would report `None`, never 100
Accounting invariant	cash + open value ≈ equity — holds at this snapshot

> **Equity is quoted at cost, and that is a deliberate limitation.** The open book is valued at $510 — its cost — because an unpriced position is counted at cost rather than at zero: counting it at zero would report a loss nobody observed. **The true mark-to-market value of the 51 open positions is not stated here**, and on this market it could be materially lower. Only the **realised** figure, –$53.51, is a settled fact.

07 The arithmetic that decides this strategy

+$2.50 + +$2.50 + +$2.50 + +$2.50 = −$10.00

FOUR clean +25% winners
…and the four only get you back to flat

ONE rug
a single position, gone entirely

A TARGET HIT IS NOT A PROFIT

The strategy's most important assumption is false in production. Proceeds are the router's number, not `1.25 × cost`. Of **53 target hits: 14 returned less than the nominal $12.50, and 2 returned less than the $10 that went in** — the smallest was **$5.08**, a 49% loss on a trade the ledger calls a win. This is the guard working as designed, and it means “hit rate” is not the same as “win rate”.

THE AVERAGE HIDES THE DISTRIBUTION

Target hits averaged **$14.085** — but the median was **$12.837** and one outlier returned **$68.73**. Strip that single trade and the picture changes materially. **A handful of outliers is currently carrying the strategy's entire positive side**, which is exactly the shape that does not survive a small sample.

CAPITAL TIED UP INDEFINITELY

There is no stop and no time exit. A position at 0.05× that the provider has not marked inactive is held **forever**, and its $10 is unavailable for anything else. Its hold reason is recorded as `below_target` — the same reason as a position at 1.24×. The wallet cannot tell you which of its 51 open slots are alive.

RISK	WHAT PRODUCTION SHOWS
Fees & slippage	Average **entry price impact 10.24%** — charged into the entry price, so the 1.25× target is measured from an already-degraded fill.
Never reaches 1.25×	Unbounded. 51 positions are open with no deadline; any of them may simply never qualify.
Dead tokens	**27 of 80** closed trades returned exactly $0.00 — a 33.8% total-loss rate against a strategy needing 71–80% winners.
Market-data gaps	Readings older than 900 s are refused. During a gap nothing can hit its target, but a token can still die — the two directions are not symmetric.
Opportunity cost	Skips are **permanent** — never revisited when cash frees up. Deliberate, so the capture rate means what it says, but a real cost.
Duplicate protection	Working: 131 → 131 → 131, zero duplicates. Also a ceiling — a token that later becomes a better opportunity can never be re-entered.
Track Record quality	Karthik inherits the Radar's admission standard wholesale and re-applies no filter, so its results are as good as admission is — and no better.

WHERE THE STRATEGY STANDS RIGHT NOW

Hit rate achieved	66.3%
Hit rate needed (realised avg)	71.0%
Hit rate needed (median)	77.9%
Hit rate needed (nominal)	80.0%
Realised P&L	−$53.51

Below every threshold, and the realised figure agrees. **Three hours is not a verdict** — but it is not encouraging either, and it should not be read as one.

The asymmetry, stated plainly. The upside of any single trade is capped at +25% by the rule itself. The downside is capped at −100% by reality. A strategy that caps its winners and not its losers needs a very high hit rate to survive, and needs it consistently. That is the whole bet.

08 What must be measured before calling this a success

Rising equity is not evidence. With no stop loss and no time exit, open positions are marked at cost until they resolve — so equity can sit flat or drift upward while the book quietly fills with tokens that will settle at $0.00. **The only settled number is realised P&L**, and it is currently **–$53.51**. Judge on the list below, on a sample large enough to mean something, and not before.

MEASURE	NOW	WHAT IT HAS TO ESTABLISH
Target hit rate	66.3%	Sustained above the break-even threshold, not merely above it once
Dead / rug rate	33.8%	Whether it is stable, or a function of the market regime it was measured in
Expectancy per trade	–$0.67	Positive, with a confidence interval that excludes zero
Profit factor	0.80	gross win $216.49 ÷ gross loss $270.00 — must exceed 1.0
Max drawdown	not yet meaningful	Peak-to-trough on realised equity, across a full cycle
Time to 1.25×	not yet reported	How fast winners resolve — this is what sets capital turnover
Time capital stays locked	51 open, no deadline	How long non-winners hold a slot before dying or qualifying
Opportunity capture	100%	Whether it survives saturation — untested, because cash has not yet run out
Avg concurrent positions	51 and climbing	The steady state, once entries and exits balance
Fees & slippage	entry impact 10.24%	Whether impact scales with size, and how much of the edge it eats

CUTS THAT HAVE TO BE RUN SEPARATELY

- **Token age at entry** — `detected_at` vs `track_record_at`, both stored, `NULL` when unknown
- **Entry liquidity** — `entry_liquidity_usd`; the thin-pool cohort is where impact does its damage
- **Entry market cap** — `entry_market_cap`
- **Track Record score at admission** — does a better admission actually trade better?
- **Daily / regime performance** — a rising market can carry a strategy with no edge

THE HONEST THRESHOLD

This strategy is a success when **expectancy per trade is positive and profit factor exceeds 1.0 over a sample large enough that a handful of outliers cannot explain it** — measured on realised trades, across more than one market regime, with fees and slippage already charged.

It is not a success because equity rose. It is not a failure because equity fell for three hours. Neither statement is yet supported by 80 closed trades.

VERIFIED AGAINST PRODUCTION

ENVIRONMENT	`production`
HOST	`vps-651730bb · 51.79.166.133 · OVH`
PUBLIC WEBSITE	`www.memescope.site — HTTP 200`
PUBLIC API	`api.memescope.site — alpha-gated`
SERVING COMMIT	`bcc5550c2baf07aea484ae97e64de2f064626031`
HOST CHECKOUT	`11fe5077aef76e086166a70ba9007b55c60f6401` `NOT built · NOT serving`
ALEMBIC HEAD	`0044_karthik_wallet — single head`
STRATEGY ID	`karthik · wallet a0da0193-4fea-4ac1-891b-44abc897dc84`
STRATEGY VERSION	`no registry entry — the rule is app/karthik/rules.py`
GENERATED	`2026-08-22 14:08 UTC`

ASSUMPTIONS & LIMITATIONS

- **Three hours of trading.** Every rate here rests on 80 closed positions taken between 11:00 and 14:08 UTC on a single day. No conclusion about the strategy is supportable from it.
- **Open positions are marked at cost**, not at market. Equity of $946.49 is therefore an upper-ish bound, not a valuation.
- Karthik has **no registry strategy id or version** — unlike the default wallet, its rule lives in code (`app/karthik/rules.py`) rather than in a versioned registry entry. It is identified here by wallet ID and serving commit.
- Expectancy and profit factor use **realised** trades only; the 51 open positions contribute nothing to either. The API could not be queried programmatically (the alpha gate rejects scripted sign-in), so figures come from production **source, container environment and database**.
- Production moved from `bcc5550` to `11fe507` on disk during verification; every statement describes the **running image**, still `bcc5550`. No production data, configuration, strategy or wallet was altered — nothing was activated, paused or reset.